

HADOOP

Hadoop Map Reduce

Introduction


- Open source project written in Java
- Large scale distributed data processing
- Based on Google's Map Reduce framework and Google file system
- Works on commodity hardware
- Used by a Google, Yahoo, Facebook, Amazon, and many other startups

<http://wiki.apache.org/hadoop/PoweredBy>

Hadoop Core

- Hadoop Distributed File System (HDFS)
 - Distributes and stores data across a cluster (brief intro only)
- Hadoop Map Reduce (MR)
 - Provides a parallel programming model
 - Moves computation to where the data is
 - Handles scheduling, fault tolerance
 - Status reporting and monitoring

Typical Cluster



- Nodes are Linux PCs
- 4-8GB RAM ~100s of GB IDE/SATA drives

GFS

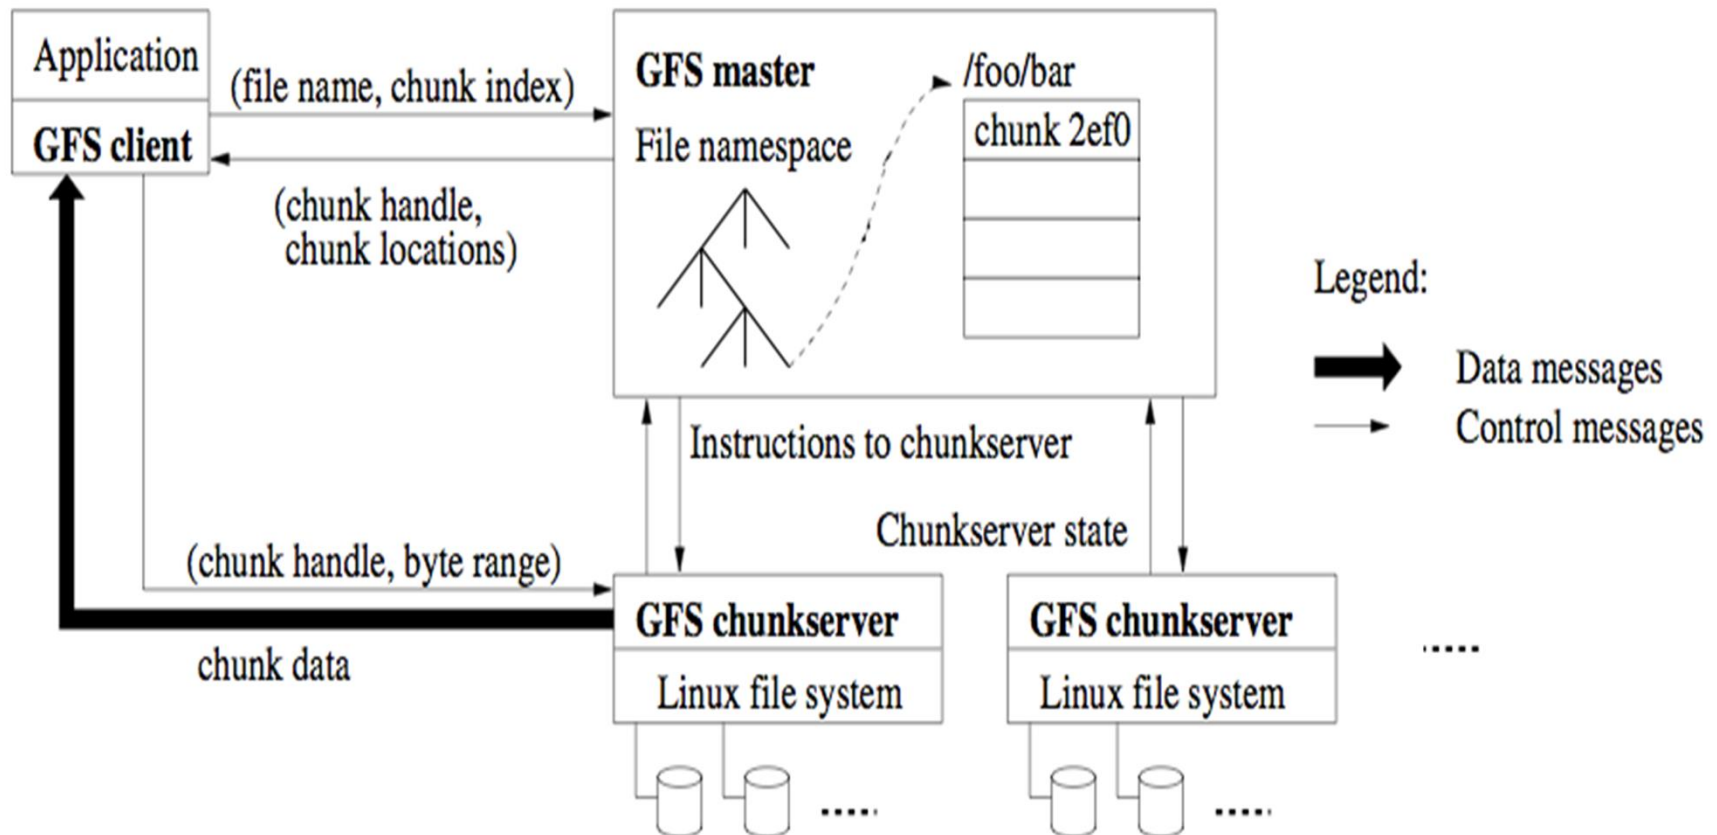
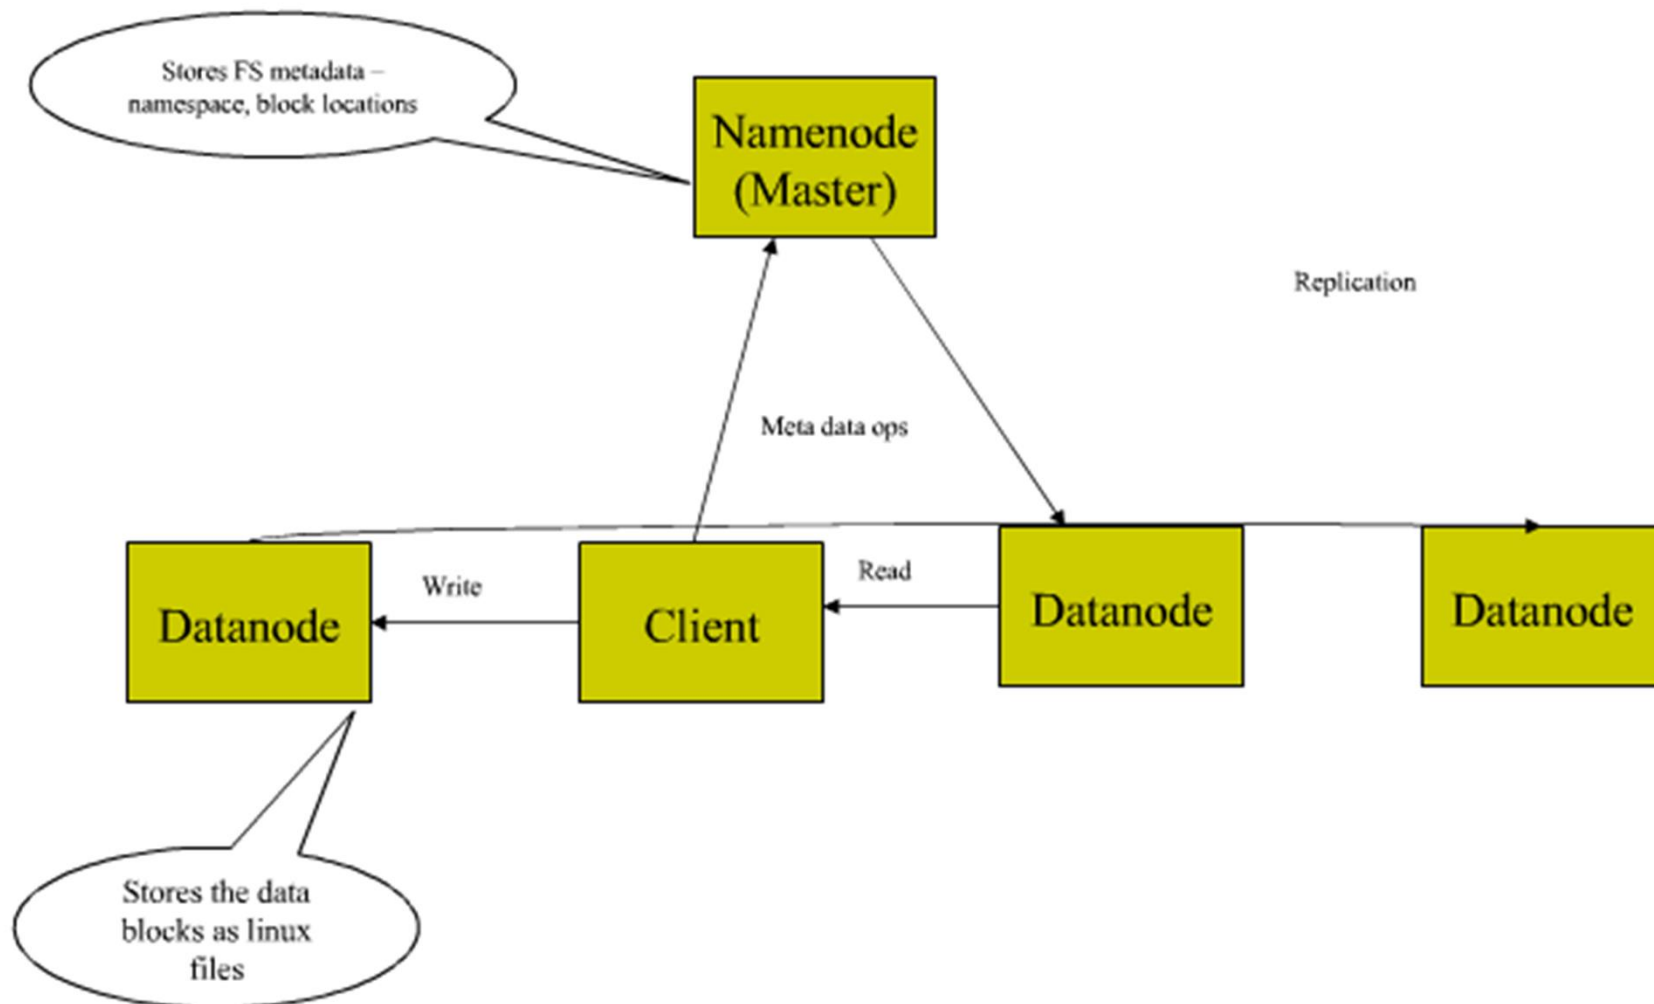


Figure 1: GFS Architecture

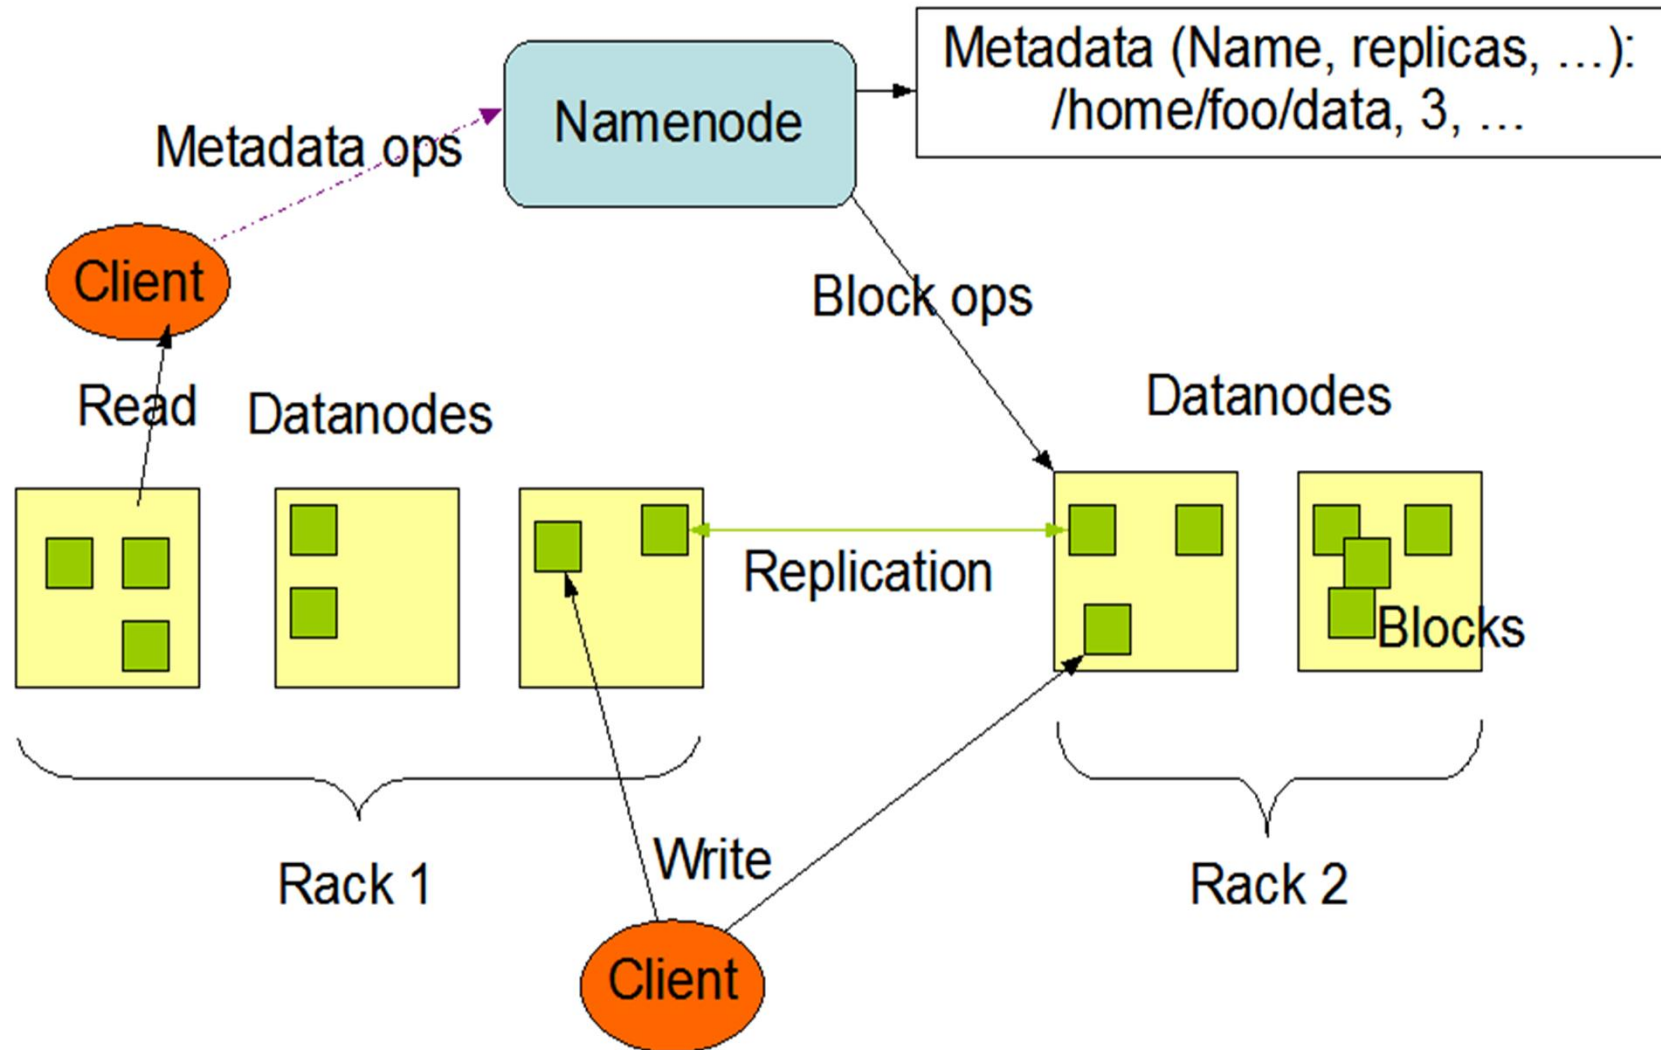
HADOOP Distributed File System


- ❑ Scale to Petabytes across 1000s of nodes
- ❑ Single namespace for entire cluster
- ❑ Files broken into 128MB blocks
- ❑ Block level replication handles node failure
- ❑ Optimized for single write multiple reads
- ❑ Writes are append only

HDFS Architecture



Hadoop Architecture





Word Count Problem

- ❑ Find the frequency of each word in a given corpus of documents
- ❑ Trivial for small data
- ❑ How to process more than a TB of data
- ❑ Doing it on one machine is very slow – takes days to finish!
- ❑ Good News : It can be parallelized across number of machines


Hadoop Map Reduce

- ❑ Why Map Reduce
- ❑ Map Reduce Architecture - 1
- ❑ Map Reduce Programming Model
- ❑ Word count using Map Reduce
- ❑ Map Reduce Architecture - 2

Why Map Reduce

- ❑ How to scale large data processing applications ?
- ❑ Divide the data and process on many nodes
- ❑ Each such application has to handle
 - Communication between nodes
 - Division and scheduling of work
 - fault tolerance
 - monitoring and reporting
- ❑ Map Reduce handles and hides all these issues
- ❑ Provides a clean abstraction for programmer

Map Reduce Architecture



- ❑ Each node is part of a HDFS cluster.
- ❑ Input data is stored in HDFS spread across nodes and replicated
- ❑ Programmer submits job (mapper, reducer, input) to Job tracker
- ❑ Job tracker - Master
 - splits input data
 - Schedules and monitors various map and reduce tasks
- ❑ Task tracker - Slaves
 - Execute map and reduce tasks

Map Reduce Programming Model

- Inspired by functional language primitives
- *map f list* : applies a given function f to a each element of list and returns a new list
$$\text{map square } [1\ 2\ 3\ 4\ 5] = [1\ 4\ 9\ 16\ 25]$$
- *reduce g list* : combines elements of list using function g to generate a new value
$$\text{reduce sum } [1\ 2\ 3\ 4\ 5] = [15]$$
- Map and reduce do not modify input data. They always create new data
- A Hadoop Map Reduce job consists of a mapper and a reducer

Map Reduce Programming Model

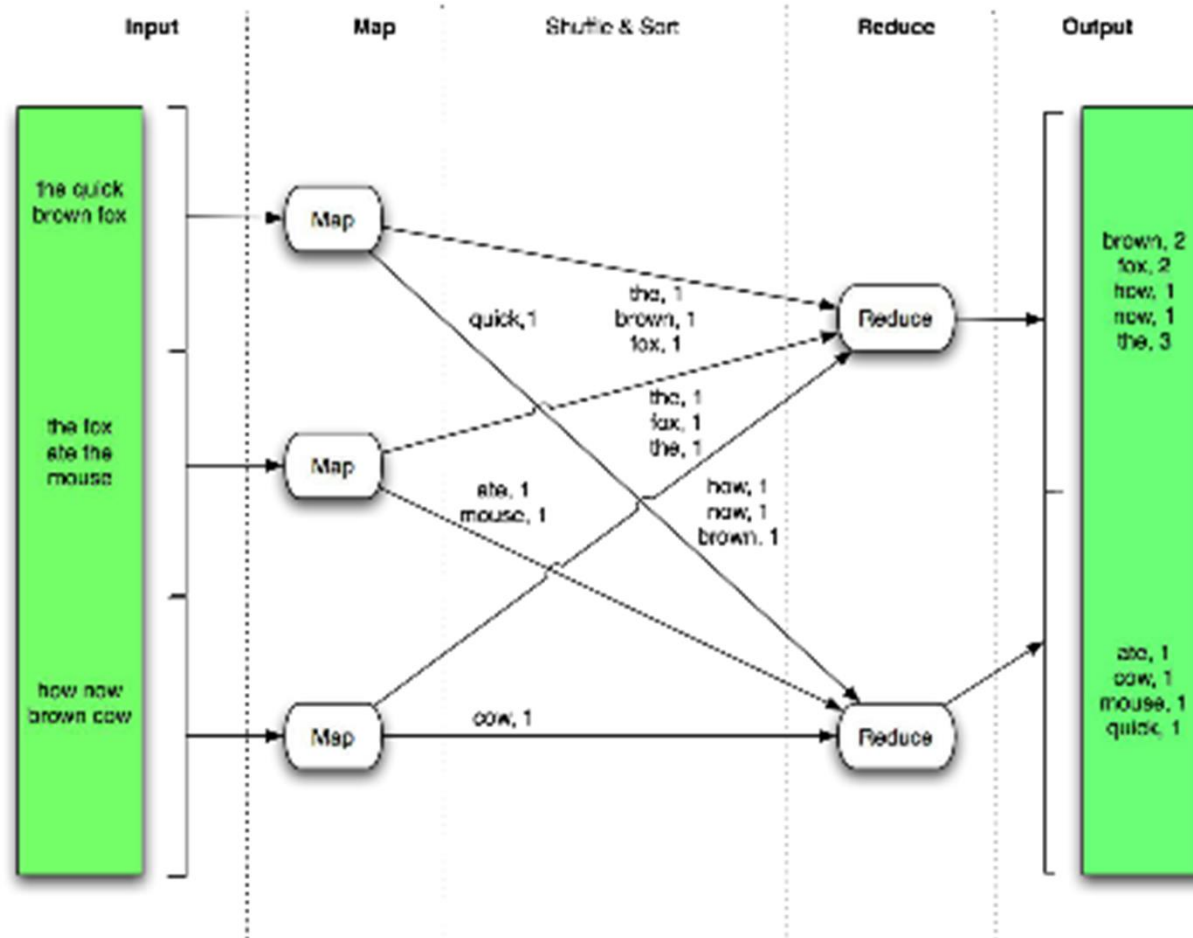
□ Mapper

- **Input:** <key:., offset, value:line of a document>
- **Output:** for each word w in input line output<key: w, value:1>
Input: (2133, *The quick brown fox jumps over the lazy dog.*)
Output: (the, 1) , (quick, 1), (brown, 1) ... (fox,1), (the, 1)

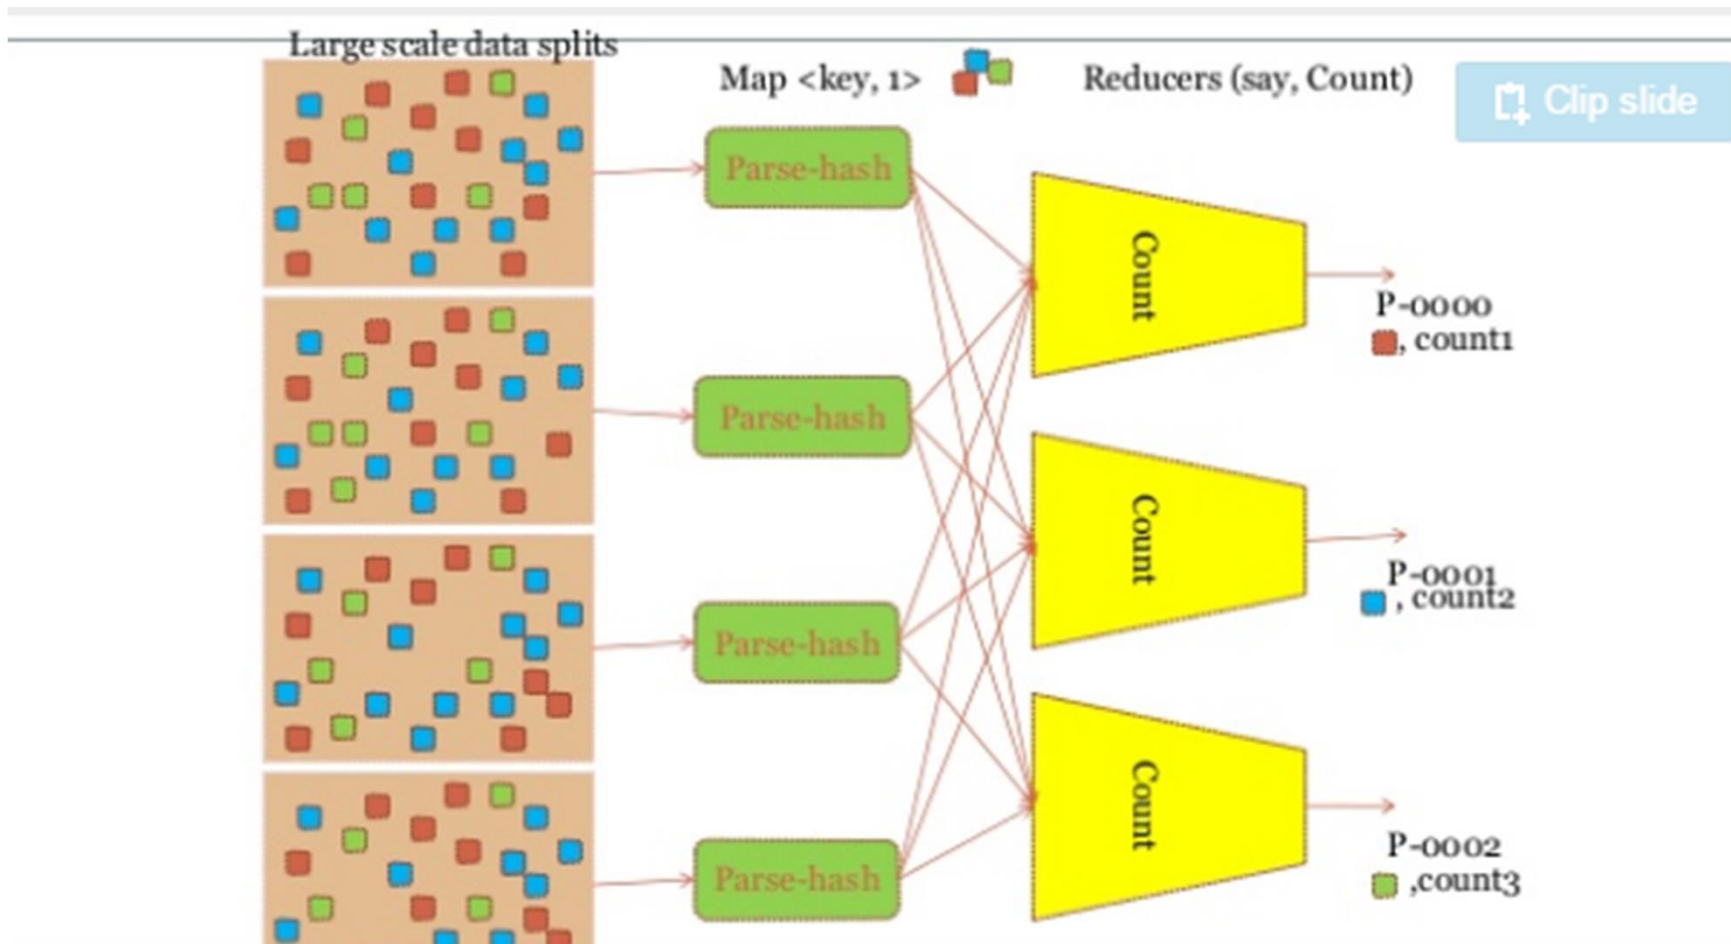
□ Reducer

- **Input:** <key: word, value: list<integer>>
- **Output:** sum all values from input for the given key input list of values
and output <Key:word value:count>
Input: (the, [1, 1, 1, 1,1]), (fox, [1, 1, 1]) ...
Output: (the, 5)
(fox, 3)

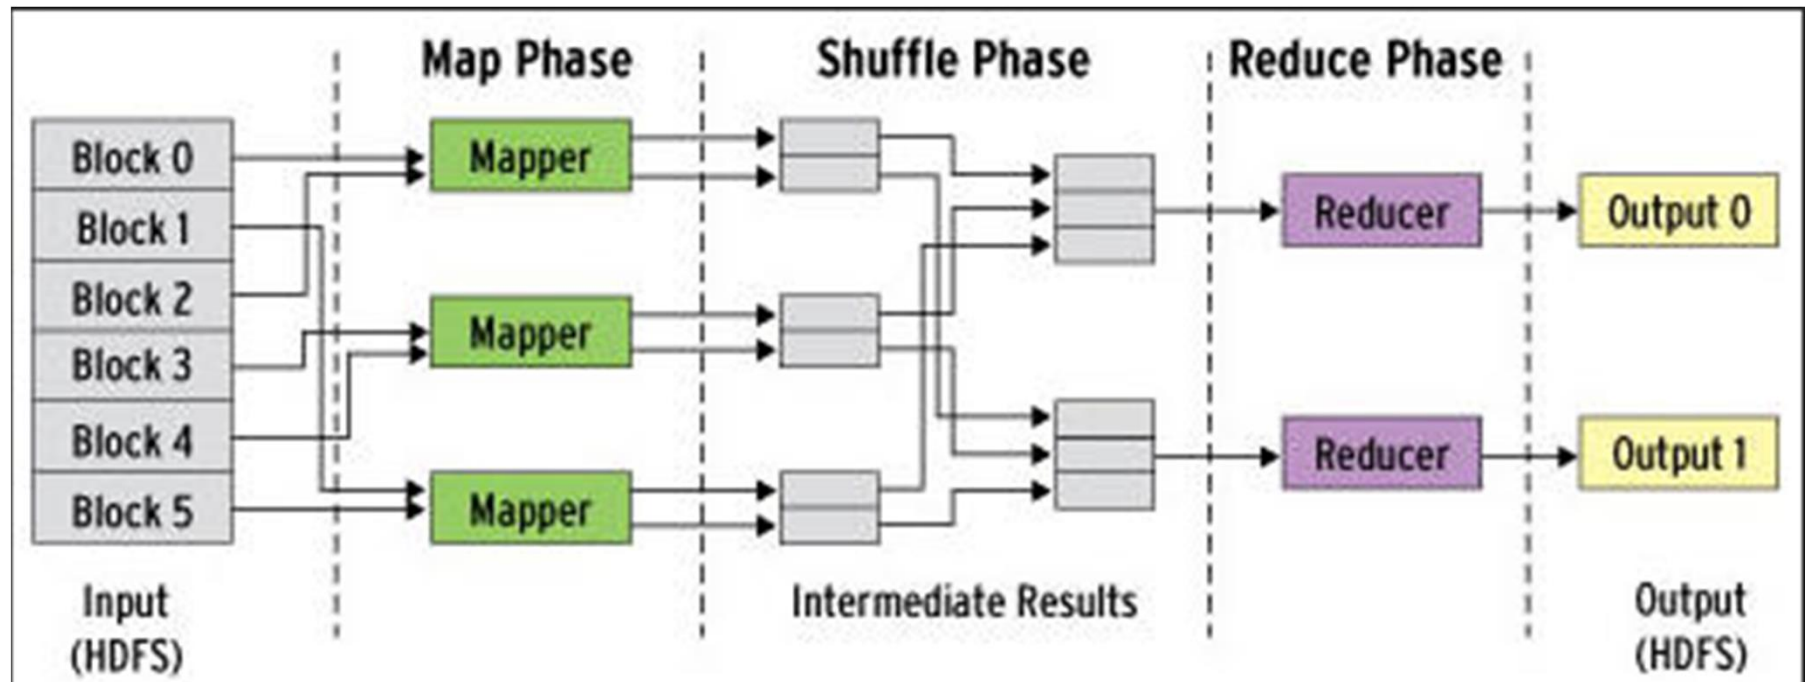
Word Count using Map Reduce



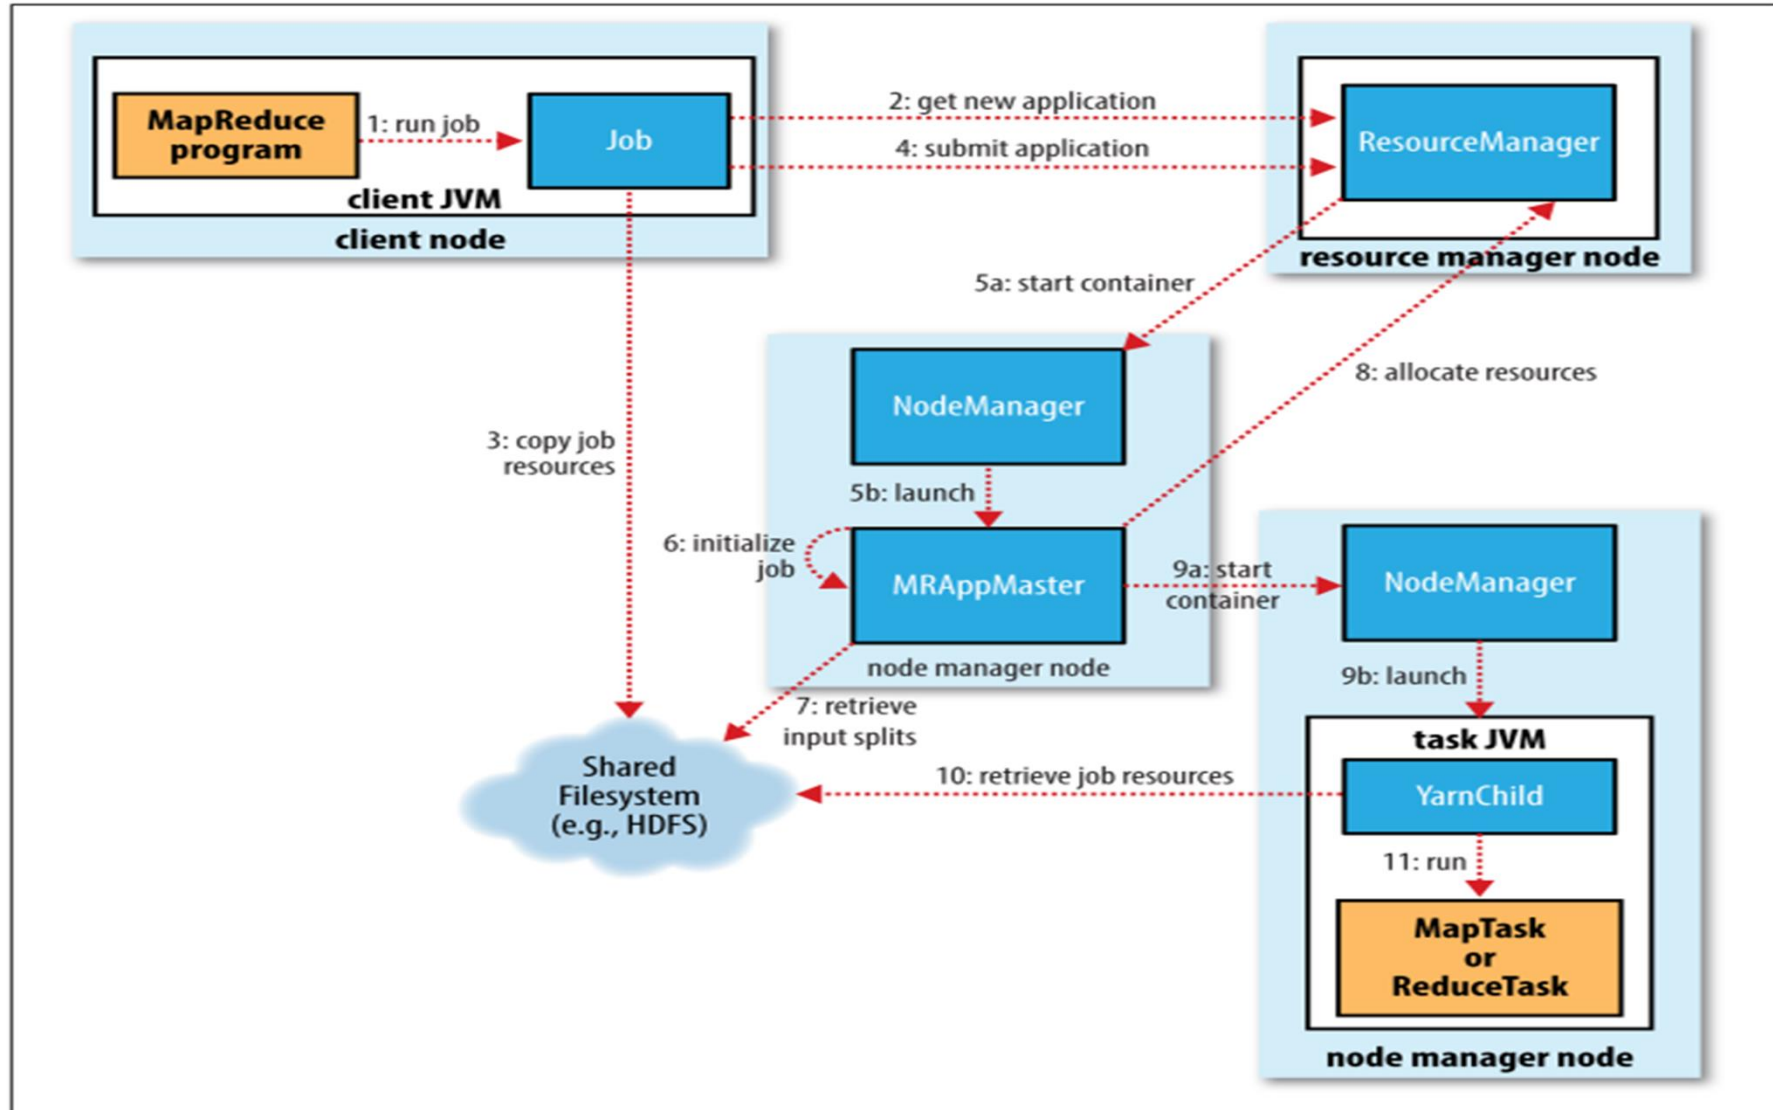
Map Reduce



Map Reduce



Map Reduce



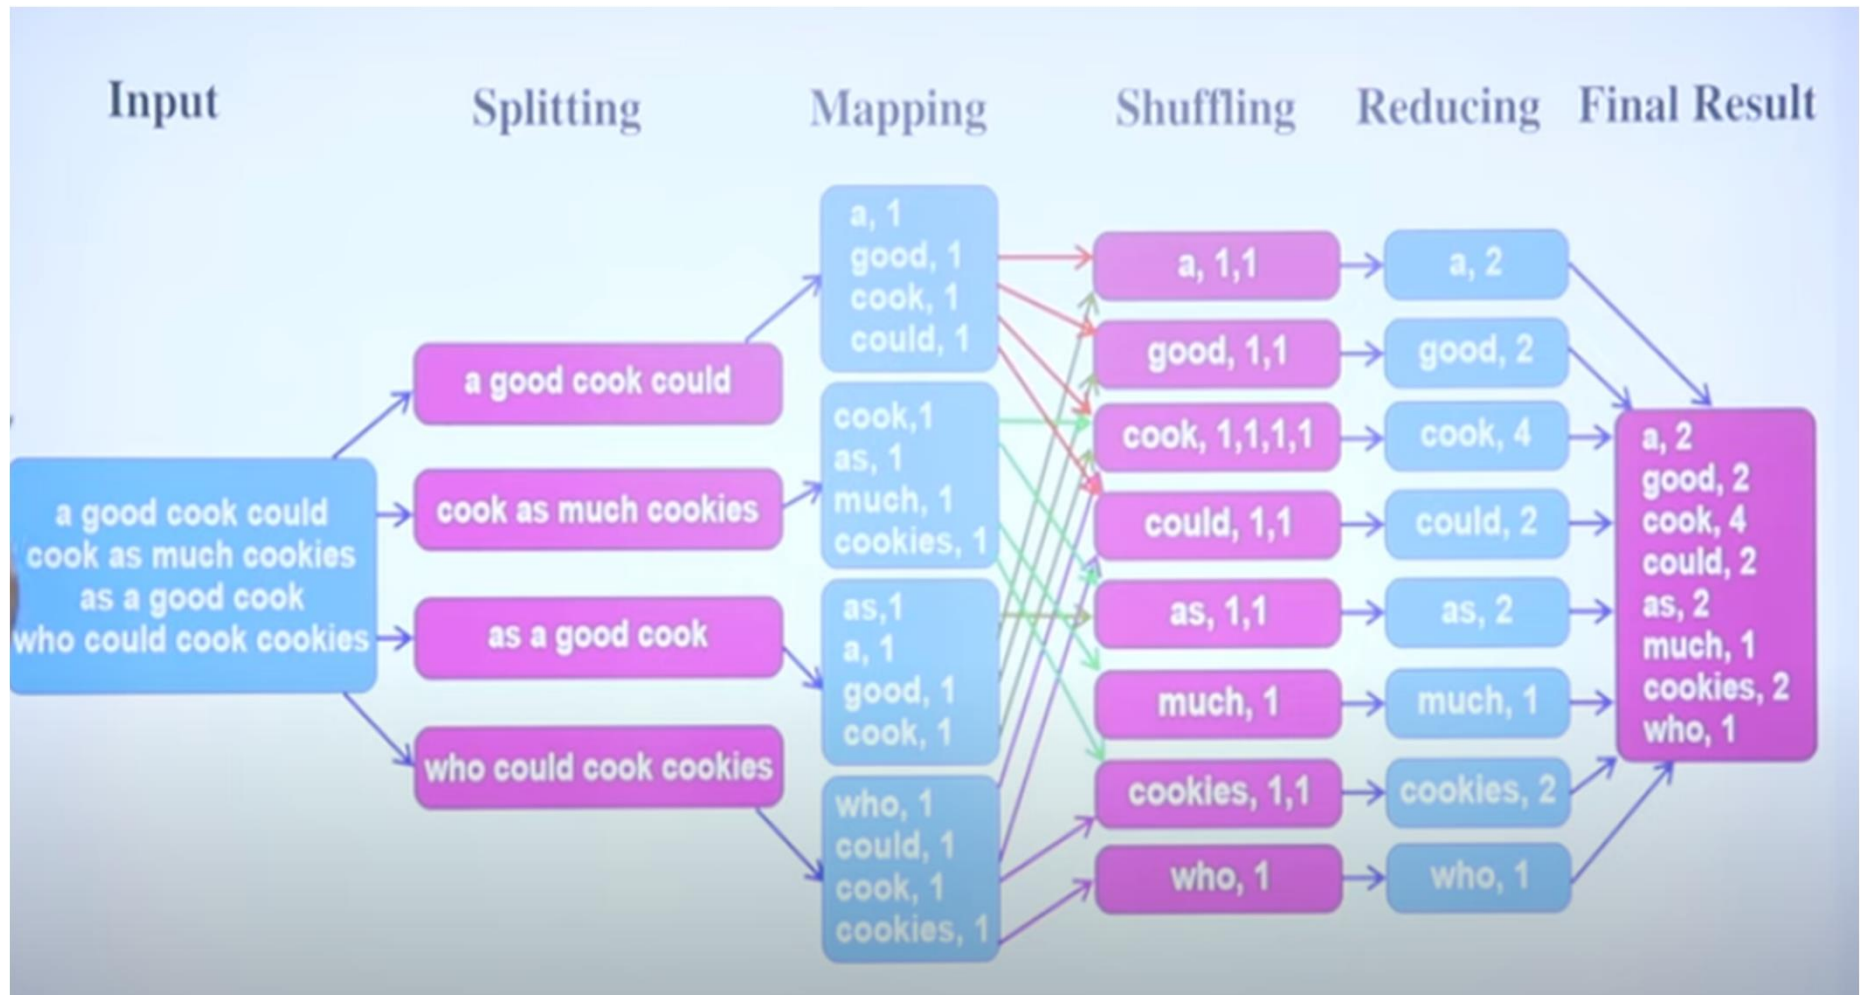
Example: Word Count

Input Files


Apple Orange Mango
Orange Grapes Plum

Apple Plum Mango
Apple Apple Plum

Why Map Reduce



Map Reduce Architecture -2



- Job tracker
 - splits input and assigns to various map tasks
 - Schedules and monitors map tasks (heartbeat)
 - On completion, schedules reduce tasks
- Task tracker
 - Execute map tasks – call mapper for every input record
 - Execute reduce tasks – call reducer for every intermediate key, list of values pair
 - Handle partitioning of map outputs
 - Handle sorting and grouping of reducer input

Map Reduce Advantages

- ❑ **Locality**
 - Job tracker divides tasks based on location of data: it tries to schedule map tasks on same machine that has the physical data
- ❑ **Parallelism**
 - Map tasks run in parallel working different input data splits
 - Reduce tasks run in parallel working on different intermediate keys
 - Reduce tasks wait until all map tasks are finished
- ❑ **Fault tolerance**
 - Job tracker maintains a heartbeat with task trackers
 - Failures are handled by re-execution
 - If a task tracker node fails then all tasks scheduled on it (completed or incomplete) are re-executed on another node